

NGINX从入门到精通进阶系列培训

实践篇：现代应用可观测

林静 Software Solution Architect

<https://myf5.net>





目录

- NGINX与Prometheus监控业务性能指标
- NGINX与EFK采集应用访问日志并可视化
- NGINX与Jaeger实现应用访问路径跟踪可视化并发现问题的实践

为什么可观测性对现代应用很重要？



什么是可观测性？

可观测性是衡量从外部输出中推断系统内部运行状态的程度 (维基百科)



现代应用的技术特征是什么？

容器化，微服务，PaaS，Cloud

你要回答的问题和你采取的措施

请求通过了哪些服务？

每个微服务在处理请求时做了什么？

如果请求很慢，瓶颈在哪里？

如果请求失败，错误发生在哪里？

请求的执行与系统的正常行为有何不同？



**很久没有接触社会了
没想到现在都这么开放了**

应用容器运行在哪里？那里发生了什么？

容器里工具不全，要啥啥没有！

这个错误日志是表达这个应用本身的问题吗？

这个服务和谁依赖，当时到底调用了哪些服务？

抓包，镜像数据？

到底是网络的问题还是应用的问题？



CNCF

Observability and Analysis

Monitoring



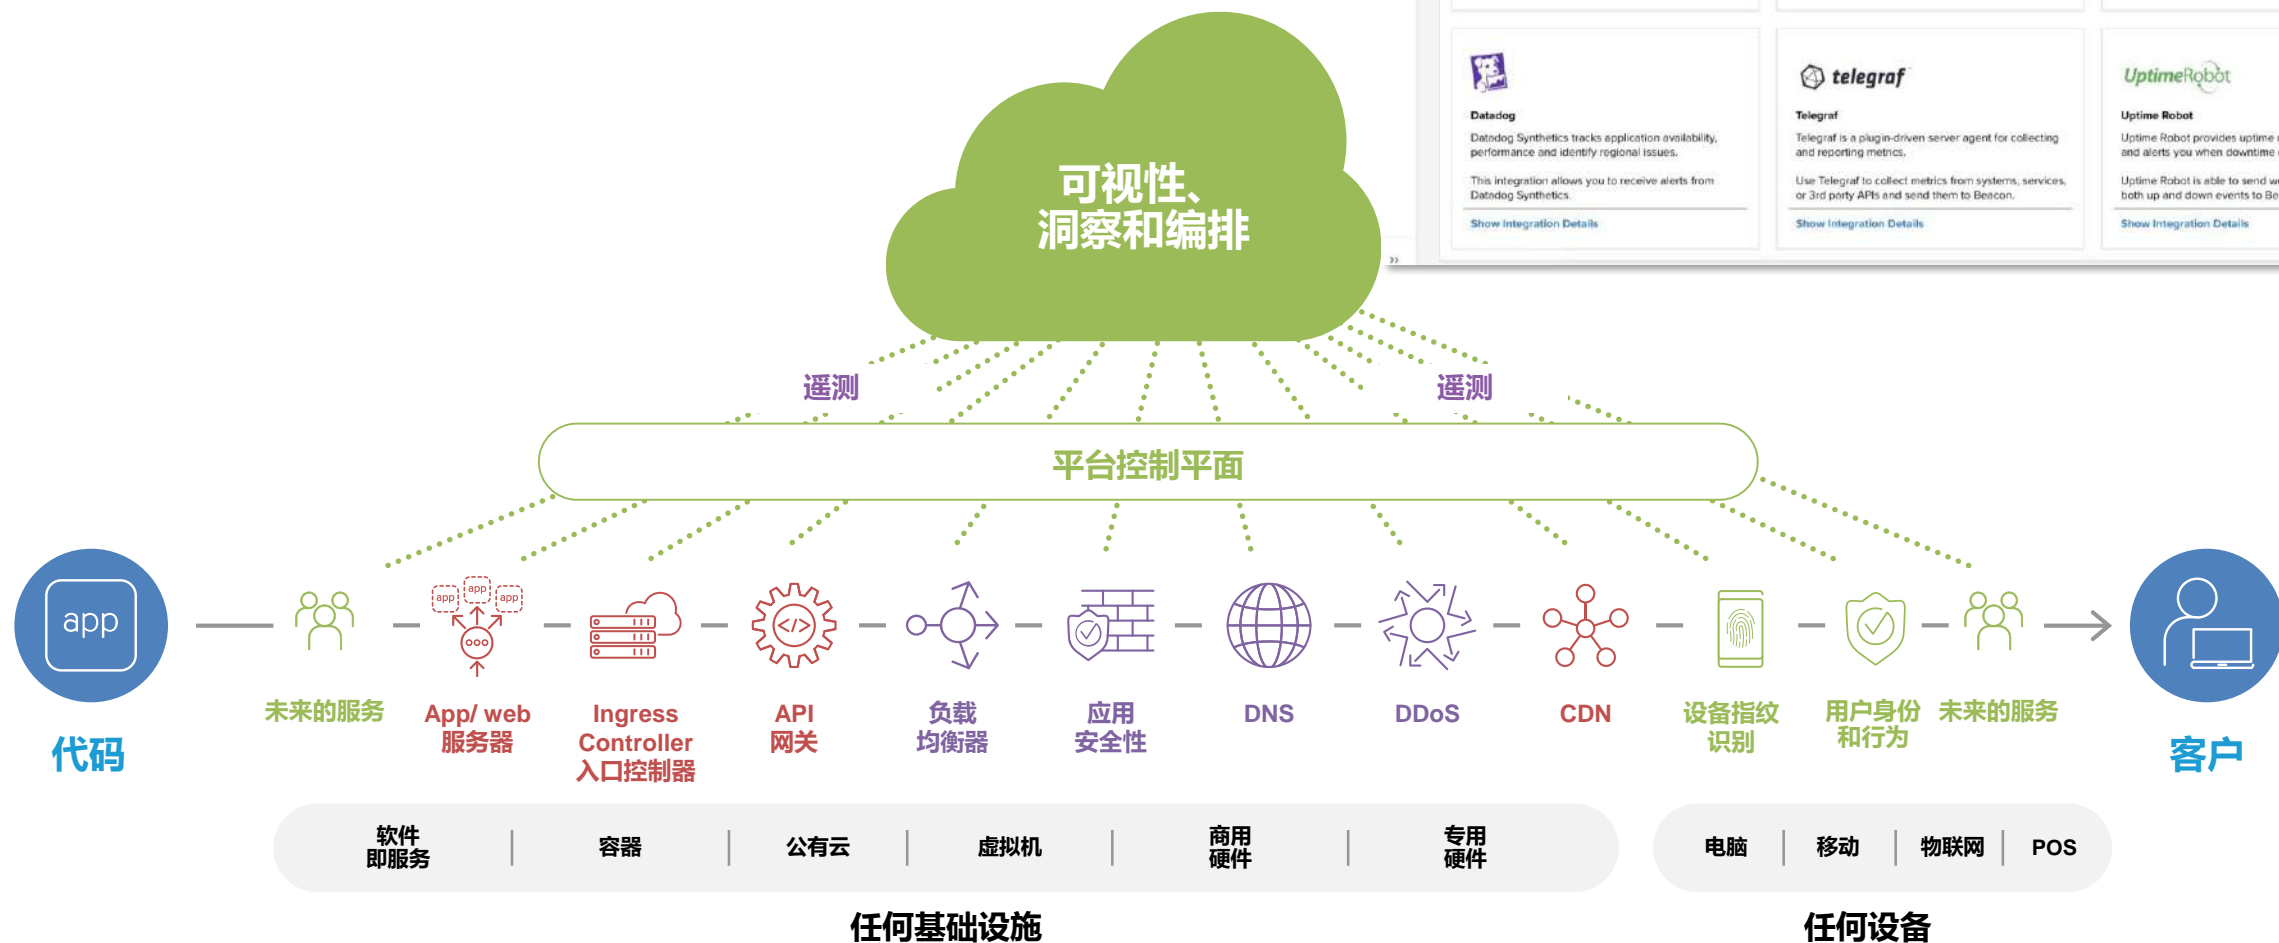
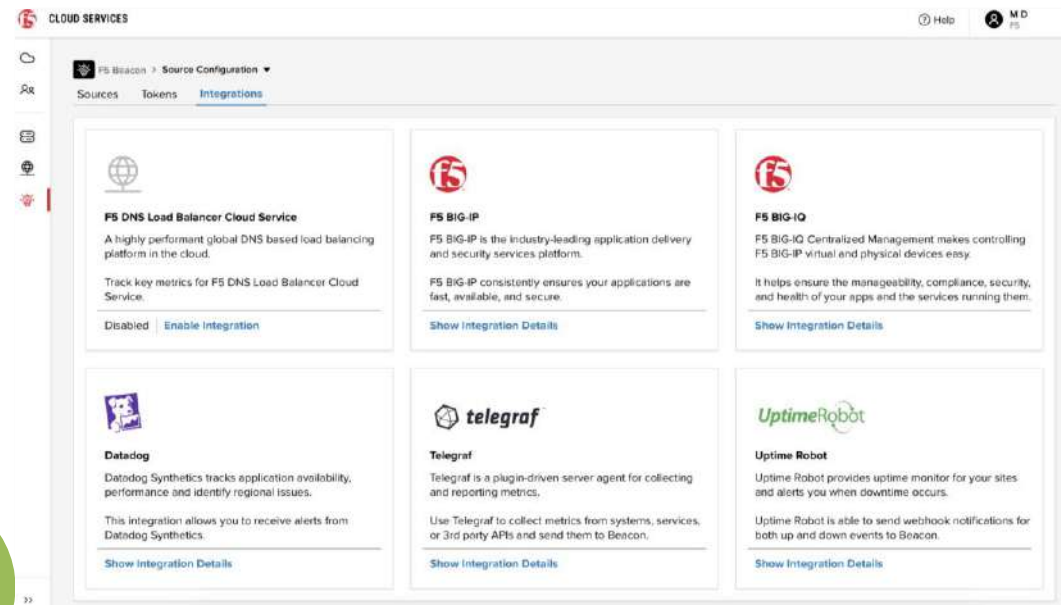
Logging

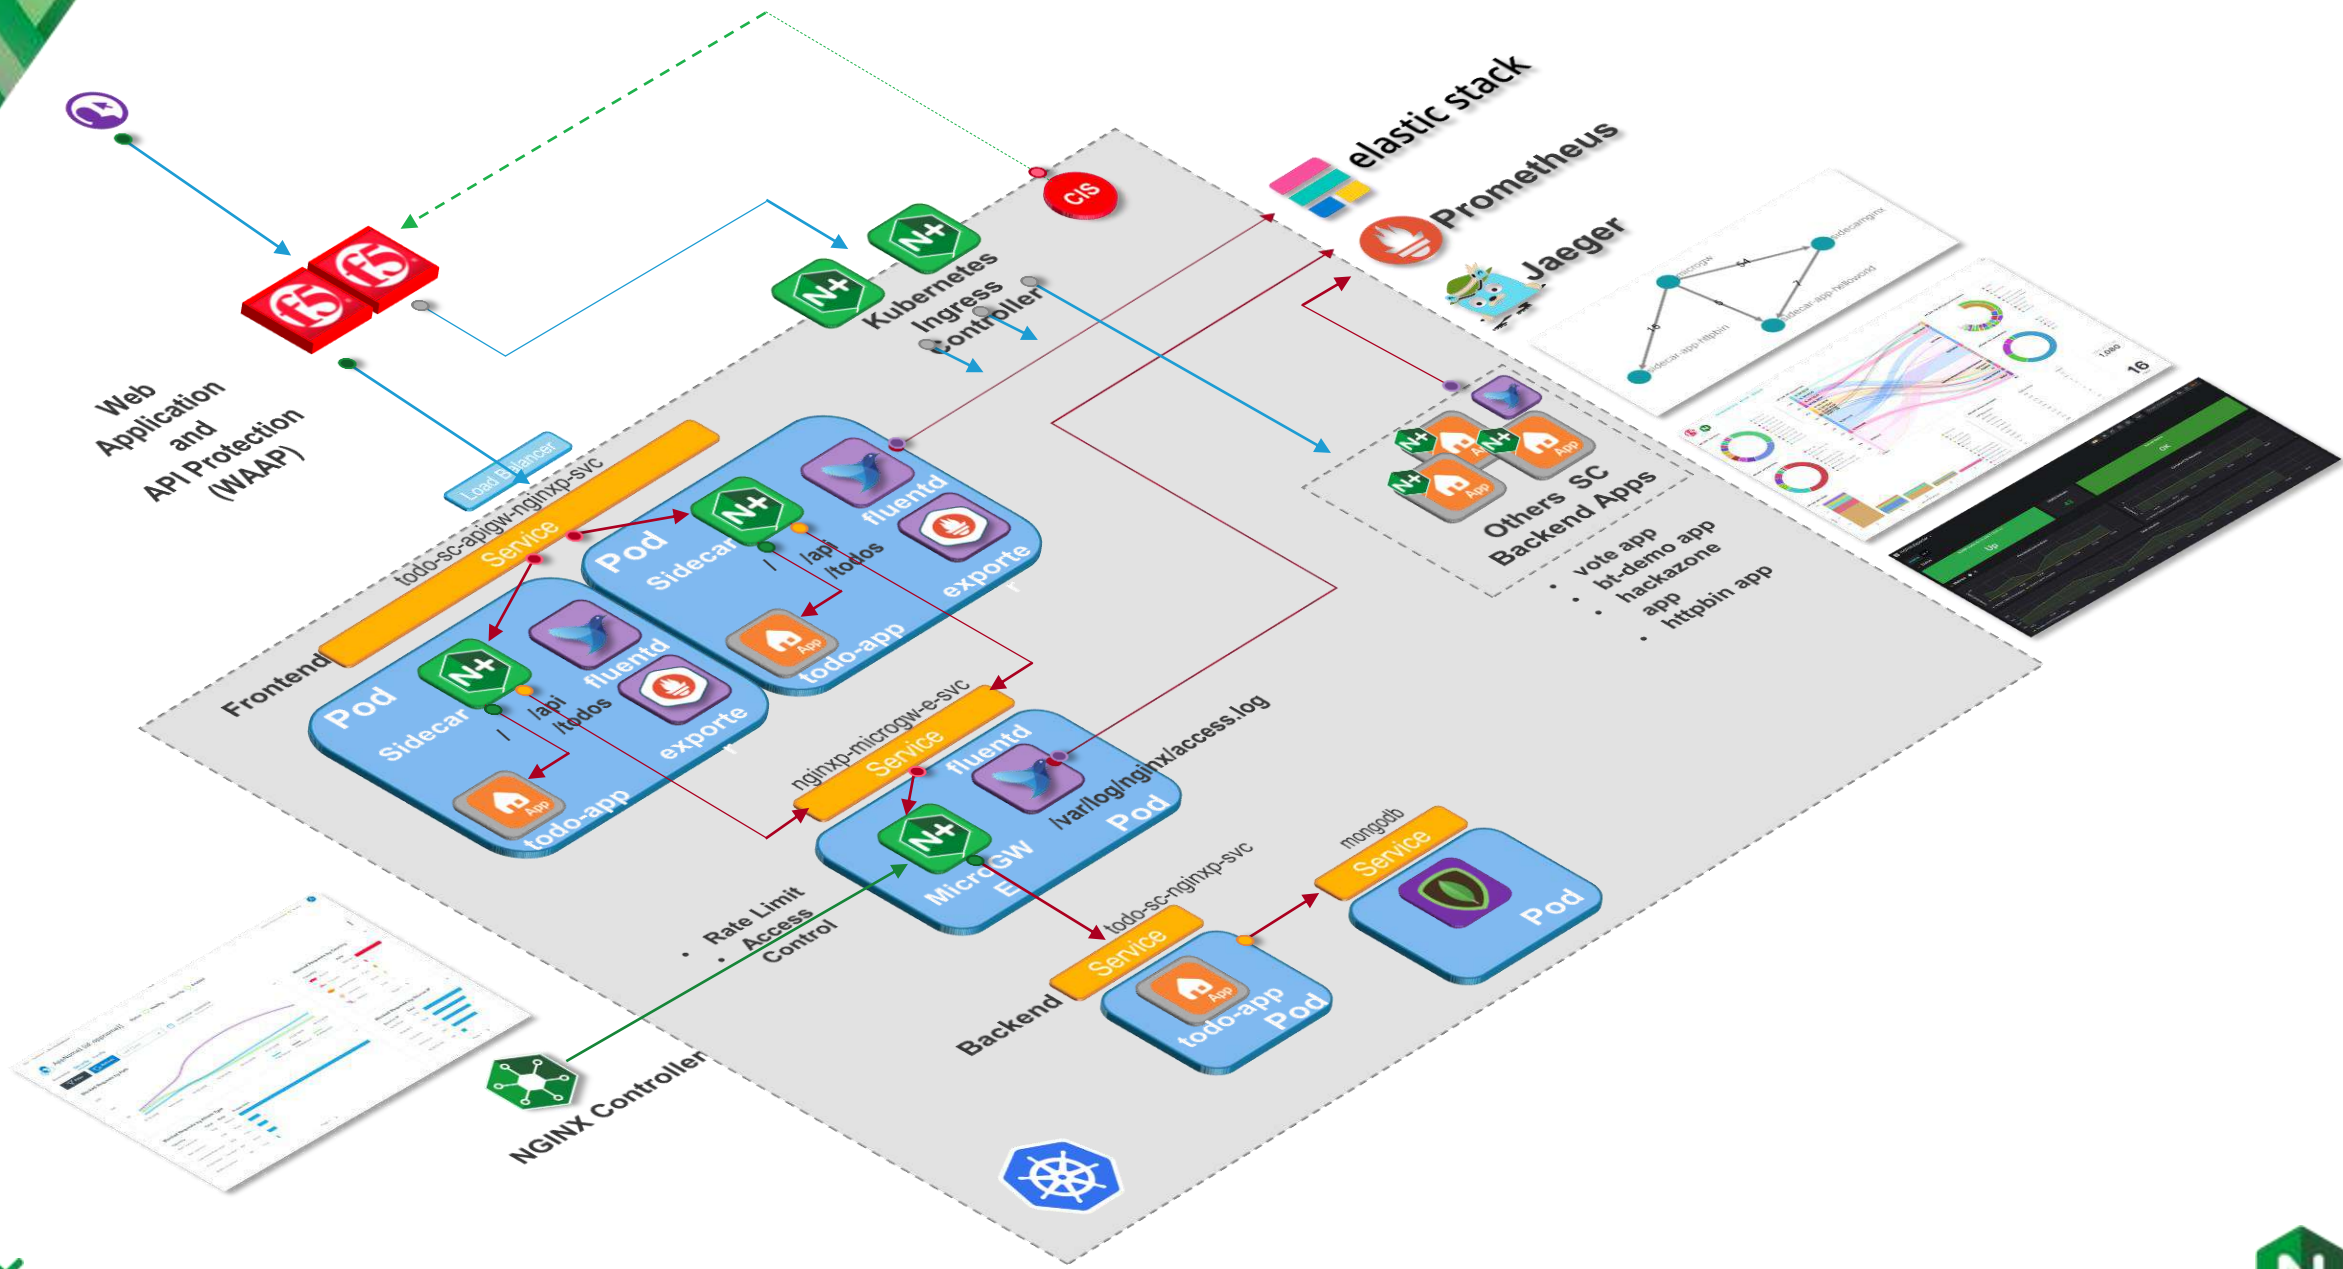


Tracing



F5在可观测领域战略



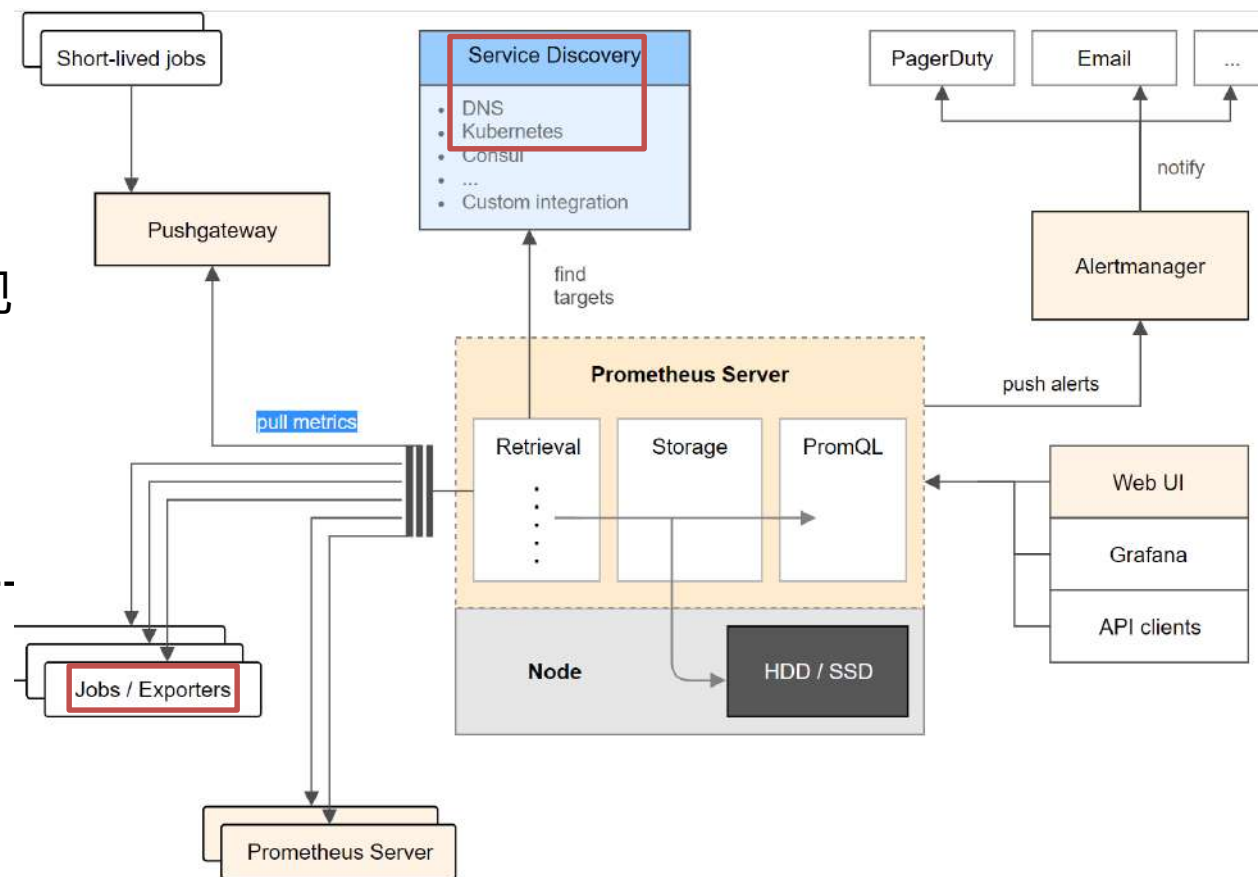


NGINX与Prometheus监控应用性能指标

NGINX

Prometheus速览

- CNCF 最“火”的项目
 - Pull 模式为主，也支持通过pushgateway做push模式
 - 目标主机通过 服务发现，静态配置或者DNS方式发现
 - 支持告警管理
 - 时序数据库，非常适合实时metric处理
 - 基于labels汇聚metrics
-
- Service discovery—config--- relabel-configs—pulling--
-metric-relabel-configs



NGINX如何与Prometheus结合?

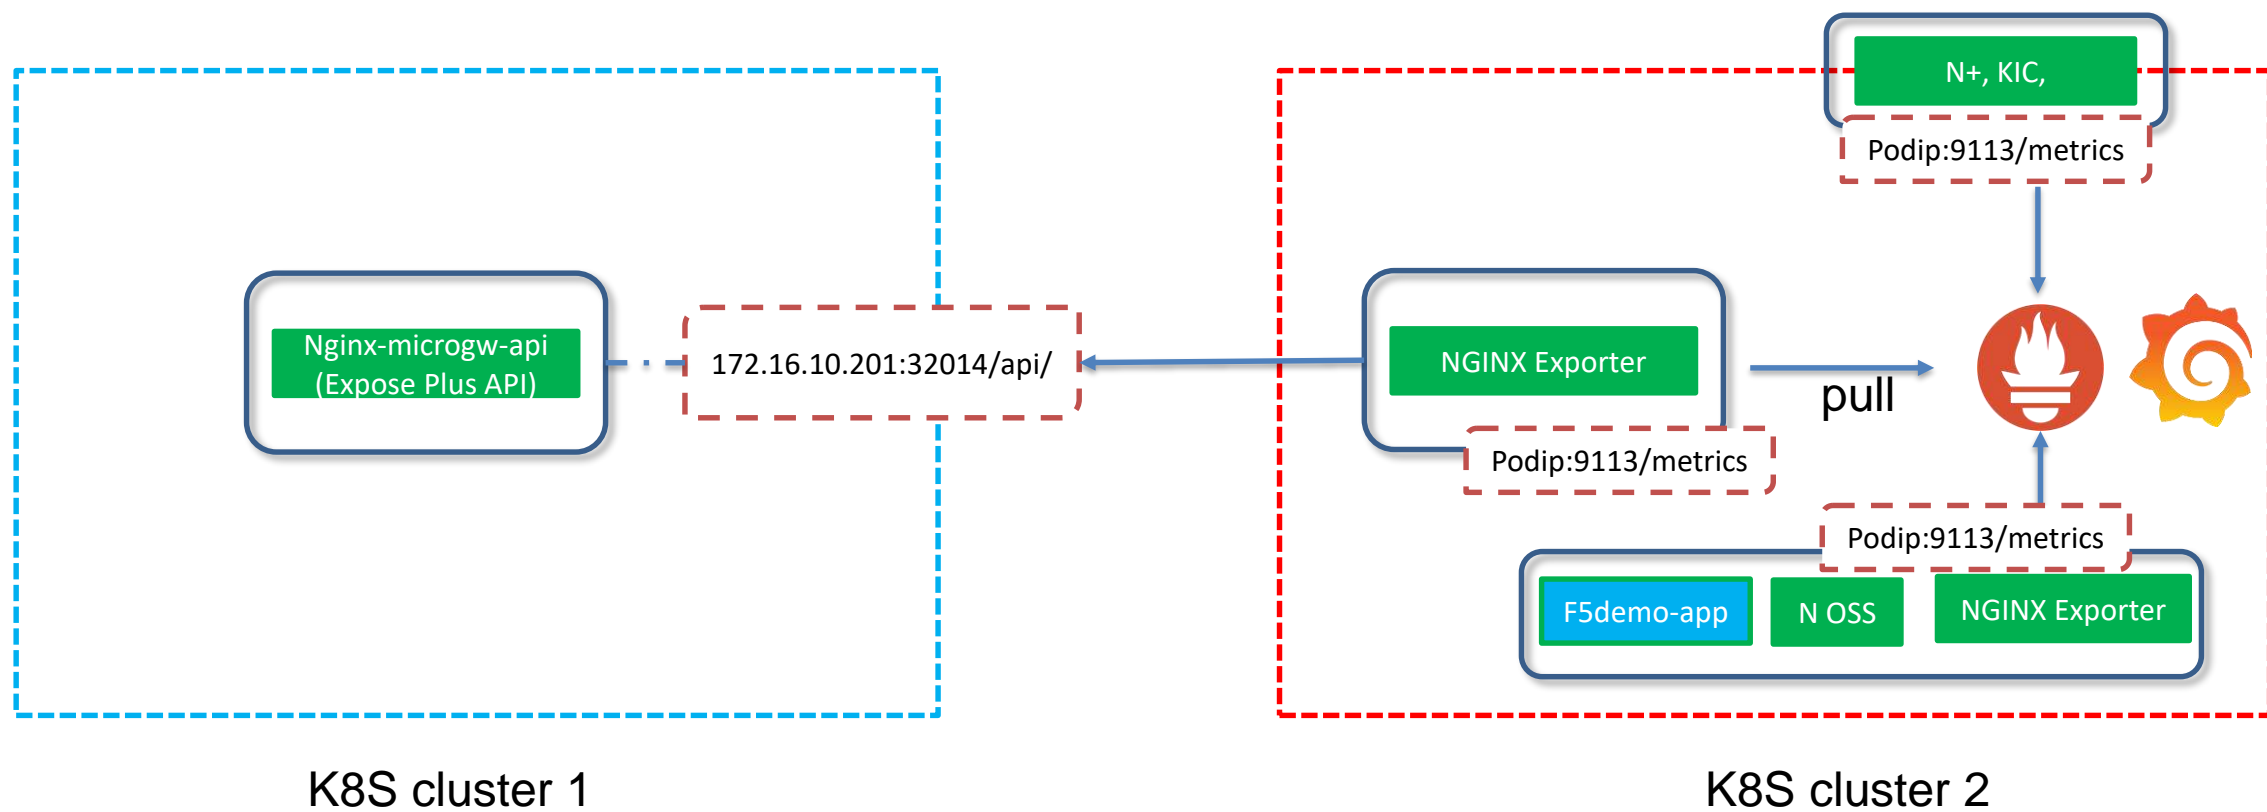
Exported Metrics

- Common metrics:
 - `nginxexporter_build_info` -- shows the exporter build information.
- For NGINX, the following metrics are exported:
 - All `stub_status` metrics.
 - `nginx_up` -- shows the status of the last metric scrape: `1` for a successful scrape and `0` for a failed one.

Connect to the `/metrics` page of the running exporter to see the complete list of metrics along with their descriptions.

- For NGINX Plus, the following metrics are exported:
 - `Connections`.
 - `HTTP`.
 - `SSL`.
 - `HTTP Server Zones`.
 - `Stream Server Zones`.
 - `HTTP Upstreams`. Note: for the `state` metric, the string values are converted to float64 using the following rule:
`"up" -> 1.0, "draining" -> 2.0, "down" -> 3.0, "unavail" -> 4.0, "checking" -> 5.0, "unhealthy" -> 6.0.`
 - `Stream Upstreams`. Note: for the `state` metric, the string values are converted to float64 using the following rule:
`"up" -> 1.0, "down" -> 3.0, "unavail" -> 4.0, "checking" -> 5.0, "unhealthy" -> 6.0.`
 - `Stream Zone Sync`.
 - `nginxplus_up` -- shows the status of the last metric scrape: `1` for a successful scrape and `0` for a failed one.
 - `Location Zones`.
 - `Resolver`.

NGINX Exporter实验环境介绍



NGINX Exporter Demo

1. 部署一个独立的nginx exporter并让其暴露另一个ks8集群1中的micgw-api NGINX Plus信息
2. 做一些流量访问模拟

<https://apidemo-kic.lab.f5se.io/gotohelloworld>

https://apidemo-kic.lab.f5se.io/#/HTTP_Methods/get_get

3. 在prometheus中查看相关metrics是否出现，并做简单的检索测试

`nginxplus_upstream_server_state`

`rate(nginxplus_upstream_server_requests[5m])`

4. 暴露k8s集群2中的Ingress controller的metrics

5. 在prometheus中查看相关ingress controller的metrics

`rate(nginx_ingress_nginxplus_http_requests_total[5m])`

6. 将nginx exporter 作为nginx的sidecar，暴露单个pod中的开源nginx信息

7. 在prometheus中查看相关开源nginx的metrics

`rate(nginx_connections_accepted[5m])`



```
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    run: ng-exporter
  name: ng-exporter
  namespace: default
spec:
  progressDeadlineSeconds: 600
  replicas: 1
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      run: ng-exporter
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
  template:
    metadata:
      creationTimestamp: null
    labels:
      run: ng-exporter
    annotations:
      prometheus.io/scrape: "true"
      prometheus.io/port: "9113"
    spec:
      containers:
      - image: nginx/nginx-prometheus-exporter:0.6.0
        imagePullPolicy: IfNotPresent
        name: ng-exporter
        command:
          - /usr/bin/exporter
          - "-nginx.plus"
          - "-nginx.scrape-uri"
          - "http://172.16.10.201:32014/api/"
        ports:
        - containerPort: 9113
          protocol: TCP
        resources: {}
        terminationMessagePath: /dev/termination-log
        terminationMessagePolicy: File
      dnsPolicy: ClusterFirst
      restartPolicy: Always
      schedulerName: default-scheduler
      securityContext: {}
      terminationGracePeriodSeconds: 30
```

NGINX Exporter Deployment

```
[root@k8s-master-v1-16 nginx-exporter]# kubectl get pod -l run=ng-exporter
```

NAME	READY	STATUS	RESTARTS	AGE
ng-exporter-68cf4b77db-4bnkn	1/1	Running	0	56m

Nginx exporter连接k8s 1集群中的microgw的API

留个问题：这样部署，有啥可以改进的地方

同时，如果你有很多nginx实例，该咋办？



Prometheus-njs

- 每个NGINX Plus暴露自己的metrics给Prometheus
- 避免expoter的 1:1关系，节省部署资源
- Only for Plus，使用API

<https://docs.nginx.com/nginx/admin-guide/dynamic-modules/prometheus-njs/>

暴露k8s-2中的IC metrics

- -enable-prometheus-metrics

```
app: nginx-ingress
template:
  metadata:
    labels:
      app: nginx-ingress
  annotations:
    prometheus.io/scrape: "true"
    prometheus.io/port: "9113"
```

```
- -nginx-plus
- -nginx-configmaps=$(POD_NAMESPACE)/nginx-config
- -default-server-tls-secret=$(POD_NAMESPACE)/default-server-secret
- -nginx-status
- -nginx-status-allow-cidrs=172.16.0.0/16,192.168.1.0/24
- -nginx-status-port=8888
#- -v=3 # Enables extensive logging. Useful for troubleshooting.
#- -report-ingress-status
#- -external-service=nginx-ingress
#- -enable-leader-election
- -enable-prometheus-metrics
- -enable-custom-resources
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: f5demo-sidecar
spec:
  replicas: 1
  selector:
    matchLabels:
      app: f5demo-sidecar
  template:
    metadata:
      labels:
        app: f5demo-sidecar
      annotations:
        prometheus.io/scrape: "true"
        prometheus.io/port: "9113"
    spec:
      volumes:
        - name: nginxdefault
          hostPath:
            path: /etc/nginx/f5devops.conf.d
      containers:
        - name: nginx-exporter-sidecar
          image: nginx/nginx-prometheus-exporter:0.6.0
          imagePullPolicy: IfNotPresent
          command:
            - /usr/bin/exporter
          ports:
            - containerPort: 9113
              protocol: TCP
        - name: nginx-sidecar
          image: nginx
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 443
            - containerPort: 8080
          volumeMounts:
            - mountPath: /etc/nginx/conf.d
              name: nginxdefault
        - name: f5demo-apps
          image: f5devcentral/f5-demo-app
          imagePullPolicy: IfNotPresent
          ports:
            - containerPort: 80
          restartPolicy: Always
```

NGINX Exporter作为sidecar

```
[root@k8s-master-v1-16 nginx-exporter]# kubectl get pod
```

NAME	READY	STATUS	RESTARTS	AGE
coffee-8c8ff9b4f-kgtwh	1/1	Running	3	95d
f5demo-sidecar-6fdcbc7fcc-72dg7	3/3	Running	1	21m

K8s-1的microgw Upstream requests

rate(nginxplus_upstream_server_requests[5m])

Execute

- insert metric at cursor -

Load time: 13
Resolution: 1s
Total time ser

Graph Console


- 1h +

1h

Until

Res. (s)

stacked



✓ {instance="10.244.2.51:9113",job="kubernetes-pods",kubernetes_namespace="default",kubernetes_pod_name="ng-exporter-68cf4b77db-4bnkn",pod_template_hash="68cf4b77db",run="ng-exporter",server="10.250.0.3:443",upstream="apidemo2backends"}
✓ {instance="10.244.2.51:9113",job="kubernetes-pods",kubernetes_namespace="default",kubernetes_pod_name="ng-exporter-68cf4b77db-4bnkn",pod_template_hash="68cf4b77db",run="ng-exporter",server="10.250.0.38:443",upstream="apidemobackends"}
✓ {instance="10.244.2.42:9113",job="kubernetes-pods",kubernetes_namespace="default",kubernetes_pod_name="ng-exporter-68cf4b77db-4bnkn",pod_template_hash="68cf4b77db",run="ng-exporter",server="10.250.0.3:443",upstream="apidemo2backends"}
✓ {instance="10.244.2.42:9113",job="kubernetes-pods",kubernetes_namespace="default",kubernetes_pod_name="ng-exporter-68cf4b77db-4bnkn",pod_template_hash="68cf4b77db",run="ng-exporter",server="10.250.0.38:443",upstream="apidemobackends"}



K8s-1的microgw Upstream状态



K8s-2 IC的total request



开源nginx 作为sidecar的指标





nginx_up

nginxexporter_build_info

nginxplus_connections_accepted

nginxplus_connections_active

nginxplus_connections_dropped

nginxplus_connections_idle

nginxplus_http_requests_current

nginxplus_http_requests_total

nginxplus_server_zone_discarded

nginxplus_server_zone_processing

nginxplus_server_zone_received

nginxplus_server_zone_requests

nginxplus_server_zone_responses

nginxplus_server_zone_sent

nginxplus_ssl_handshakes

nginxplus_ssl_handshakes_failed

nginxplus_ssl_session_reuses

nginxplus_up

nginxplus_upstream_keepalives

nginxplus_upstream_server_active

nginxplus_upstream_server_fails

nginxplus_upstream_server_header_time

nginxplus_upstream_server_received

nginxplus_upstream_server_requests

nginxplus_upstream_server_response_time

nginxplus_upstream_server_responses

nginxplus_upstream_server_sent

nginxplus_upstream_server_state

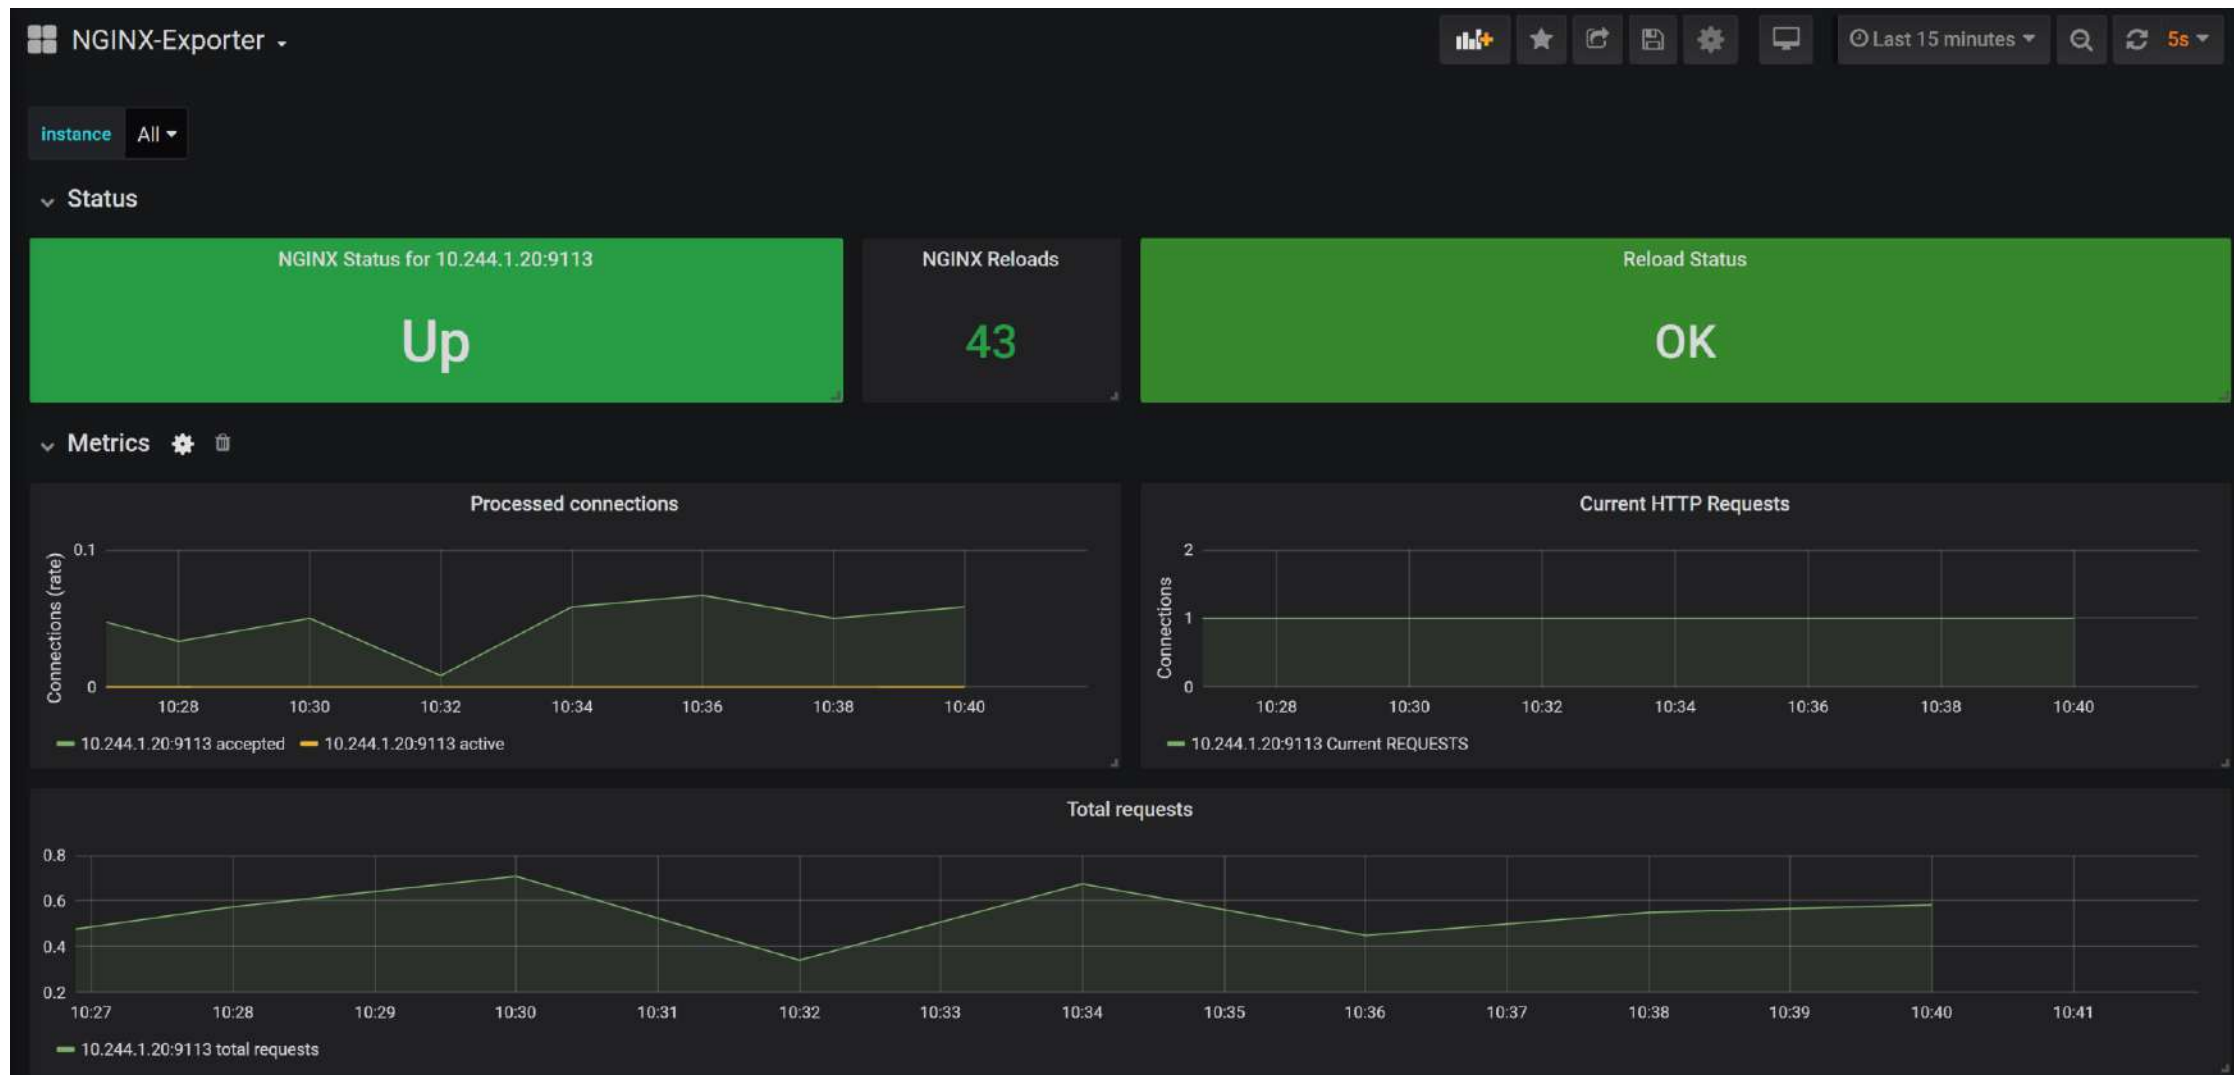
nginxplus_upstream_server_unavail

nginxplus_upstream_zombies

Exporter输出的部分指标

<https://github.com/nginxinc/nginx-prometheus-exporter>

开源NGINX指标的Grafana dashboard



结合EFK通过NGINX采集业务访问日志

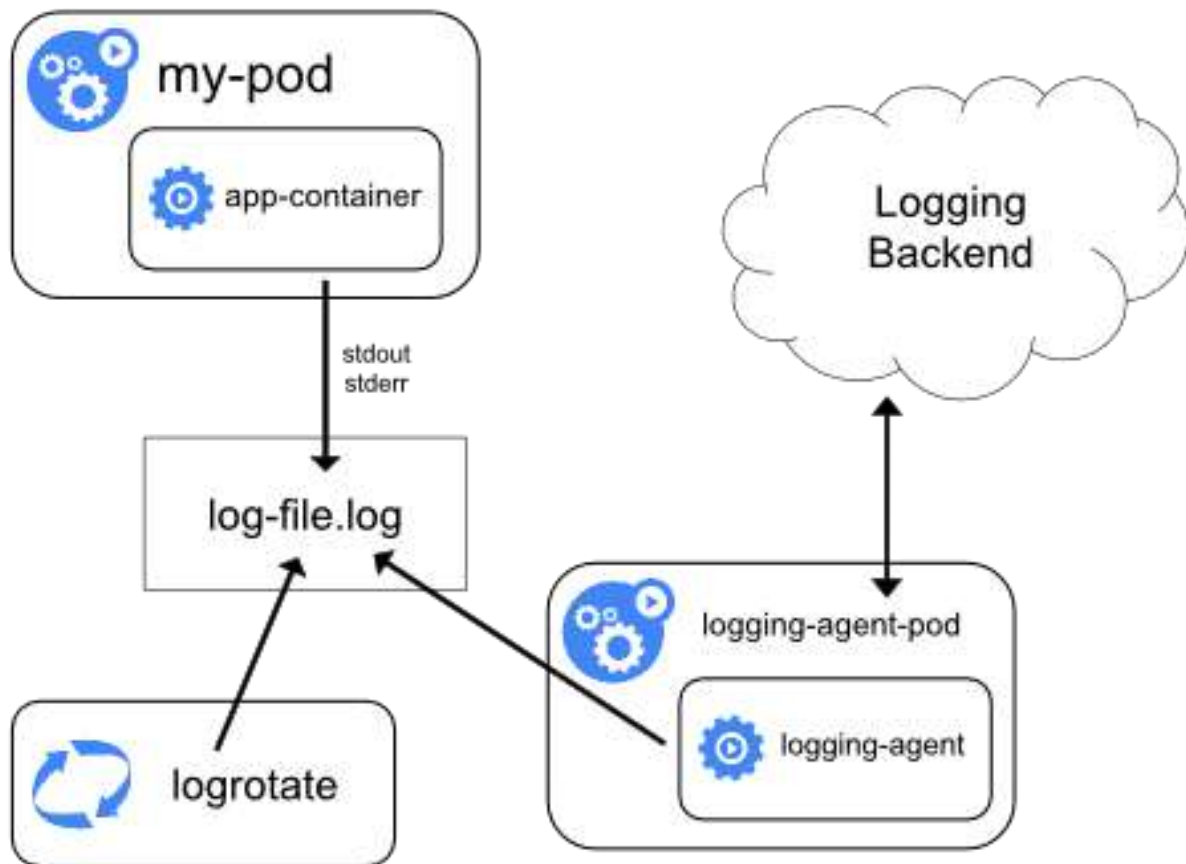
NGINX



一个问题

为什么要采集业务访问日志？？ 请抢答！

K8s下日志的采集方式-节点采集



- K8s node节点上安装日志采集程序
- 默认情况下, dockers以json-file驱动来存储容器日志到宿主机目录:
- `/var/lib/docker/containers/**containersID**/**.json`
- 日志采集程序例如fluentd采集宿主机上文件
- 采集方案难度以及灵活性相对适中
- *应用日志需要输出到stdout/stderr



节点日志采集方式demo

1. 部署EFK套件，fluentd采集日志，elasticsearch负责存储，Kibana负责可视化
2. Kibana简单检索日志，并形成可视化
3. Fluentd 在node节点采集nginx的access log
4. 当NGINX作为容器化方式使用，如何使得日志发送到stdout/stderr

部署EFK

```
[root@k8s-master-v1-16 logging]# ll
total 40
-rw-r--r-- 1 root root 2711 Dec 9 08:00 es-stateful-set.yaml
-rw-r--r-- 1 root root 382 Dec 7 22:29 es-svc.yaml
-rw-r--r-- 1 root root 16094 Dec 7 22:30 fluentd-es-configmap.yaml
-rw-r--r-- 1 root root 2438 Dec 7 22:31 fluentd-es-ds.yaml
-rw-r--r-- 1 root root 1144 Dec 9 21:04 kibana-deploy.yaml
-rw-r--r-- 1 root root 472 Dec 9 08:19 kibana-ingress.yaml
-rw-r--r-- 1 root root 354 Dec 9 08:11 kibana-svc.yaml
```

elasticsearch-logging-0	1/1	Running	0	77d
fluentd-es-v2.8.0-2fc2x	1/1	Running	4	77d
fluentd-es-v2.8.0-qkpdq	1/1	Running	0	77d
kibana-logging-57bcb74977-fd7nx	1/1	Running	2	77d

<https://github.com/kubernetes/kubernetes/tree/master/cluster/addons/fluentd-elasticsearch>

Kibana检索与展示



Time → _source

Feb 25, 2020 @ 13:12:13.508

```
kubernetes.container_name: nginx-plus-ingress kubernetes.namespace_name: nginx-ingress kubernetes.pod_name: nginx-ingress-7b67474576-ggqdl kubernetes.container_image: myf5/nginx-plus-ingress-opentracing:edge kubernetes.container_image_id: docker-pullable://myf5/nginx-plus-ingress-opentracing@sha256:9001afc1794997ce1313e86ef387add5b4fa3ccbee11456fd17b31377dd482b9 kubernetes.labels.app: nginx-ingress stream: stdout docker.container_id: 60a63488821f244b8623a62e8744f3e21464bd28910ae4c642fff72c423d2d36 kubernetes.pod_id: b2c98c58-ba82-44df-af53-1a7fdf543ad8 kubernetes.host: k8s-node1-v1-16.lab.f5se.io kubernetes.labels.pod-template-hash: 7b67474576 kubernetes.master_url: https://10.96.0.1:443/api kubernetes.namespace_id: f08e664f-990e-475a-82e6-55a8a3e6d23d message: 192.168.1.254 - - [25/Feb/2020:05:12:12 +0000] "GET /api/v1/label/.__name__/values?_=1582607532677 HTTP/1.1" 200 2215 "http://prom.lab.f5se.io/graph"
```

Expanded document

[View surrounding documents](#) [View single document](#)

Table JSON

@timestamp	Feb 25, 2020 @ 13:12:13.508
_id	vFDcEnABDPVohI5kidiy
_index	logstash-2020.02.25
_score	-
_type	_doc
docker.container_id	60a63488821f244b8623a62e8744f3e21464bd28910ae4c642fff72c423d2d36
kubernetes.container_image	myf5/nginx-plus-ingress-opentracing:edge
kubernetes.container_image_id	docker-pullable://myf5/nginx-plus-ingress-opentracing@sha256:9001afc1794997ce1313e86ef387add5b4fa3ccbee11456fd17b31377dd482b9
kubernetes.container_name	nginx-plus-ingress
kubernetes.host	k8s-node1-v1-16.lab.f5se.io
kubernetes.labels.app	nginx-ingress
kubernetes.labels.pod-template-hash	7b67474576
kubernetes.master_url	https://10.96.0.1:443/api
kubernetes.namespace_id	f08e664f-990e-475a-82e6-55a8a3e6d23d
kubernetes.namespace_name	nginx-ingress
kubernetes.pod_id	b2c98c58-ba82-44df-af53-1a7fdf543ad8
kubernetes.pod_name	nginx-ingress-7b67474576-ggqdl
message	192.168.1.254 - - [25/Feb/2020:05:12:12 +0000] "GET /api/v1/label/.__name__/values?_=1582607532677 HTTP/1.1" 200 2215 "http://prom.lab.f5se.io/graph" Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/79.0.3945.130 Safari/537.36 "-"
stream	stdout
tag	kubernetes.var.log.containers.nginx-ingress-7b67474576-ggqdl_nginx-ingress_nginx-plus-ingress-60a63488821f244b8623a62e8744f3e21464bd28910ae4c642fff72c423d2d36.log



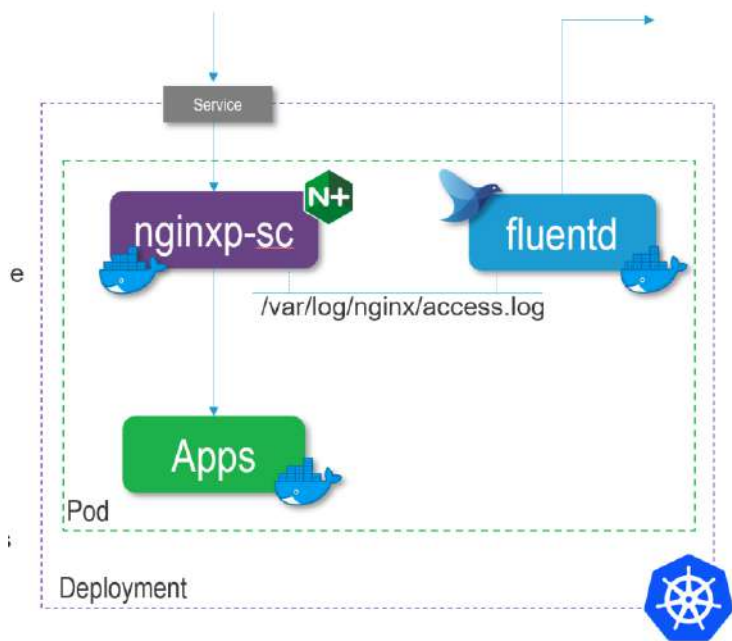

一个问题

怎么让nginx容器的access log日志输出到stdout?

制作nginx容器镜像时候，在dockerfile中：

```
RUN ln -sf /proc/1/fd/1 /var/log/nginx/access.log \  
    && ln -sf /proc/1/fd/1 /var/log/nginx/stream-access.log \  
    && ln -sf /proc/1/fd/2 /var/log/nginx/error.log
```

K8s下日志的采集方式-sidecar采集



- 当业务日志只能输出到容器内的文件时采用
- 在pod中运行一个专门的日志采集容器
- 两个容器共享相同的文件
- 采集配置自定义灵活性较高，但是技术难度及成本相对较大

Sidecar 采集方式

NGINX 和fluentd 共用同一个本地文件mount

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: sc-fluentd-td-configmap
data:
  td-agent.conf: |
    <source>
      @type tail
      format ltsv
      path /var/log/nginx/access.log
      tag nginx.access
      pos_file /var/log/td-agent/nginx-access.log.pos
      time_format %d/%b/%Y:%H:%M:%S %z
    </source>

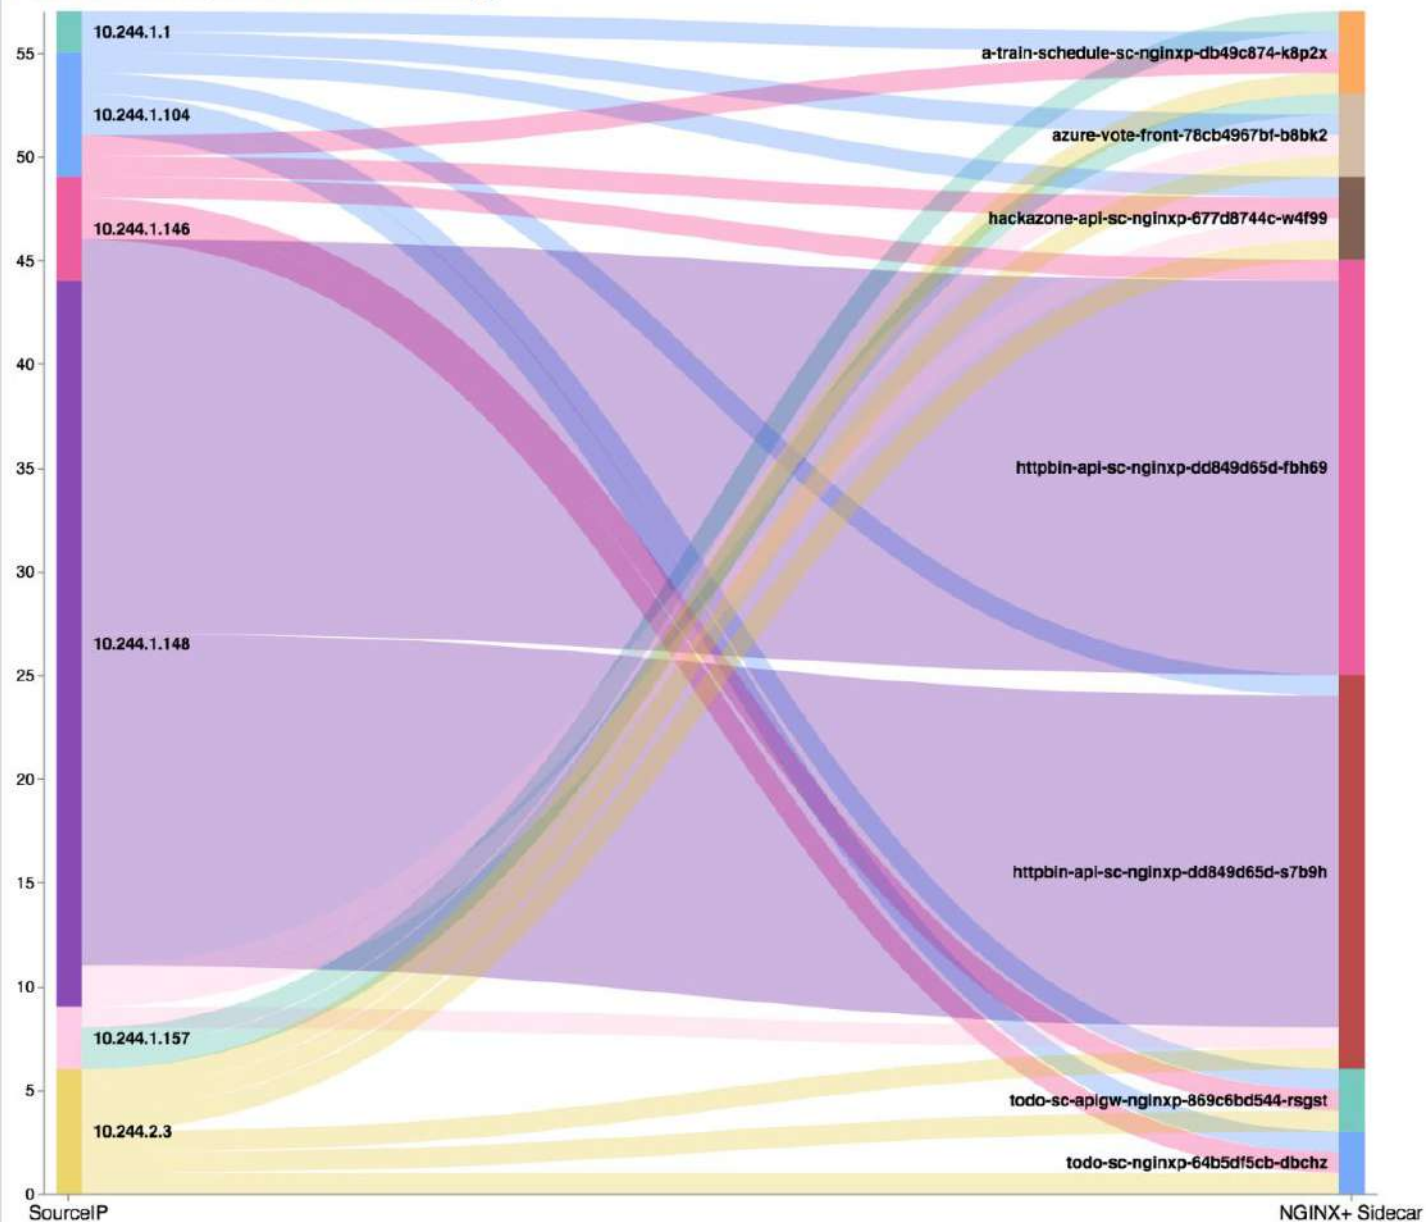
    <match nginx.*>
      @type elasticsearch
      logstash_format true
      host 192.168.211.11 # elasticsearch IP
      port 9200
      <buffer tag>
        @type memory # or file
        flush_thread_count 4
      </buffer>
      reconnect_on_error true
      reload_on_failure true
      reload_connections false
      type_name fluentd-nginx
      logstash_prefix nginxp-sidecar
    </match>
```

```
log_format ltsv 'time:$time_local\t'
                'status:$status\t'
                'request_time:$request_time\t'
                'upstream_addr:$upstream_addr\t'
                'upstream_response_time:$upstream_response_time\t'
                'upstream_cache_status:$upstream_cache_status\t'
                'body_bytes_sent:$body_bytes_sent\t'
                'remote_addr:$remote_addr\t'
                'host:$host\t'
                'hostname:$hostname\t'
                'request_method:$request_method\t'
                'request_uri:$request_uri\t'
                'protocol:$server_protocol\t'
                'x-forwarded-for: $proxy_add_x_forwarded_for\t'
                'http_referer:$http_referer\t'
                'http_user_agent:$http_user_agent';
```


```
- name: nginx-sidecar
  image: myf5/nginxp-sidecar:slim
  imagePullPolicy: Always
  ports:
    - containerPort: 443
  volumeMounts:
    - name: secret-volume
      mountPath: /etc/nginx/ssl
    - name: nginx-logs
      mountPath: /var/log/nginx
    - name: nginx-conf
      mountPath: /etc/nginx/nginx.conf
      subPath: nginx.conf
- name: fluentd
  image: reg.foobz.com.au/foobz/fluentd
  imagePullPolicy: IfNotPresent
  env:
    - name: FLUTNTD_ARGS
      value: -c /etc/td-agent/td-agent.conf
  volumeMounts:
    - name: nginx-logs
      mountPath: /var/log/nginx
    - name: config-volume
      mountPath: /etc/td-agent
  volumes:
    - name: nginx-logs
      emptyDir: {}
    - name: config-volume
      configMap:
        name: sc-fluentd-td-configmap
    - name: secret-volume
      secret:
        secretName: foobz-tls
    - name: nginx-conf
      configMap:
        name: sc-nginx-conf-tls-configmap
```



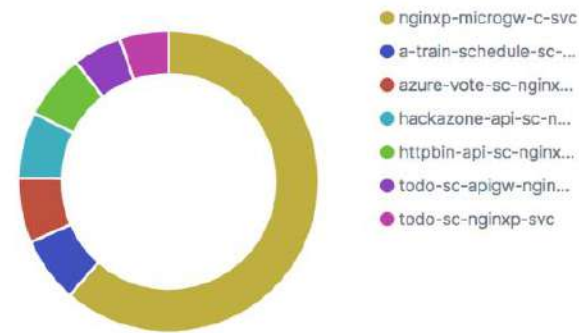
SIDECAR: Access Flow (Source IP to NGINX+ Sidecar Proxy)




SIDECAR: Top Source IP Address



SIDECAR: Popularity



SIDECAR: Top Host with Status Code



Kibana展示 <http://bde.f5se.io>



F5大数据分析展示平台
让流量数据尽收眼底，让业务态势了然于心！

抗击疫情，我们在行

客户端及请求分析

应用及响应分析

TCP层数据延迟分析

请求与响应视图-给定时段

DNS解析视图展示

系统自检测

+ Add filter


vis.服务器握手延迟(per server)




vis.L4服务器实时握手延迟



vis.服务器平均首包响应时间(per server)




vis.应用实时首包响应时间



vis.服务器连续响应平均延迟采样(每10包)



vis.应用实时服务器连续响应延迟采样




vis.首包延迟热力图



NGINX与Jaeger实现应用访问路径跟踪 可视化

NGINX


为什么要搞tracing





大家用了哪些方法去做访问路径跟踪?

还记得\$request_id吗

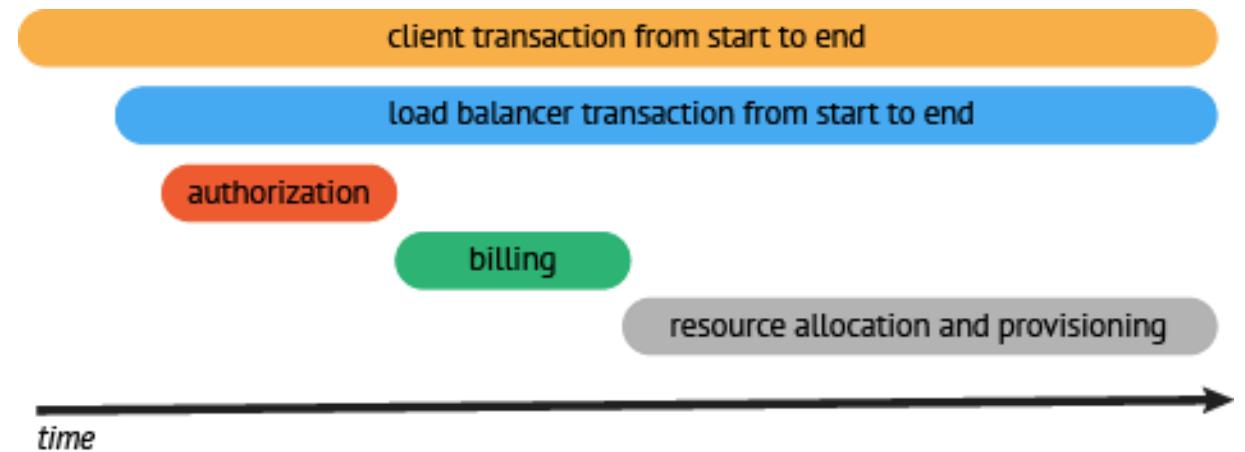



```
log_format trace '$remote_addr - $remote_user [$time_local] "$request" '
                  '$status $body_bytes_sent "$http_referer" "$http_user_agent" '
                  '"$http_x_forwarded_for" $request_id';

upstream app_server {
    server 10.0.0.1;
}

server {
    listen 80;
    add_header X-Request-ID $request_id; # Return to client
    location / {
        proxy_pass http://app_server;
        proxy_set_header X-Request-ID $request_id; # Pass to app server
        access_log /var/log/nginx/access_trace.log trace; # Log $request_id
    }
}
```

Opentracing



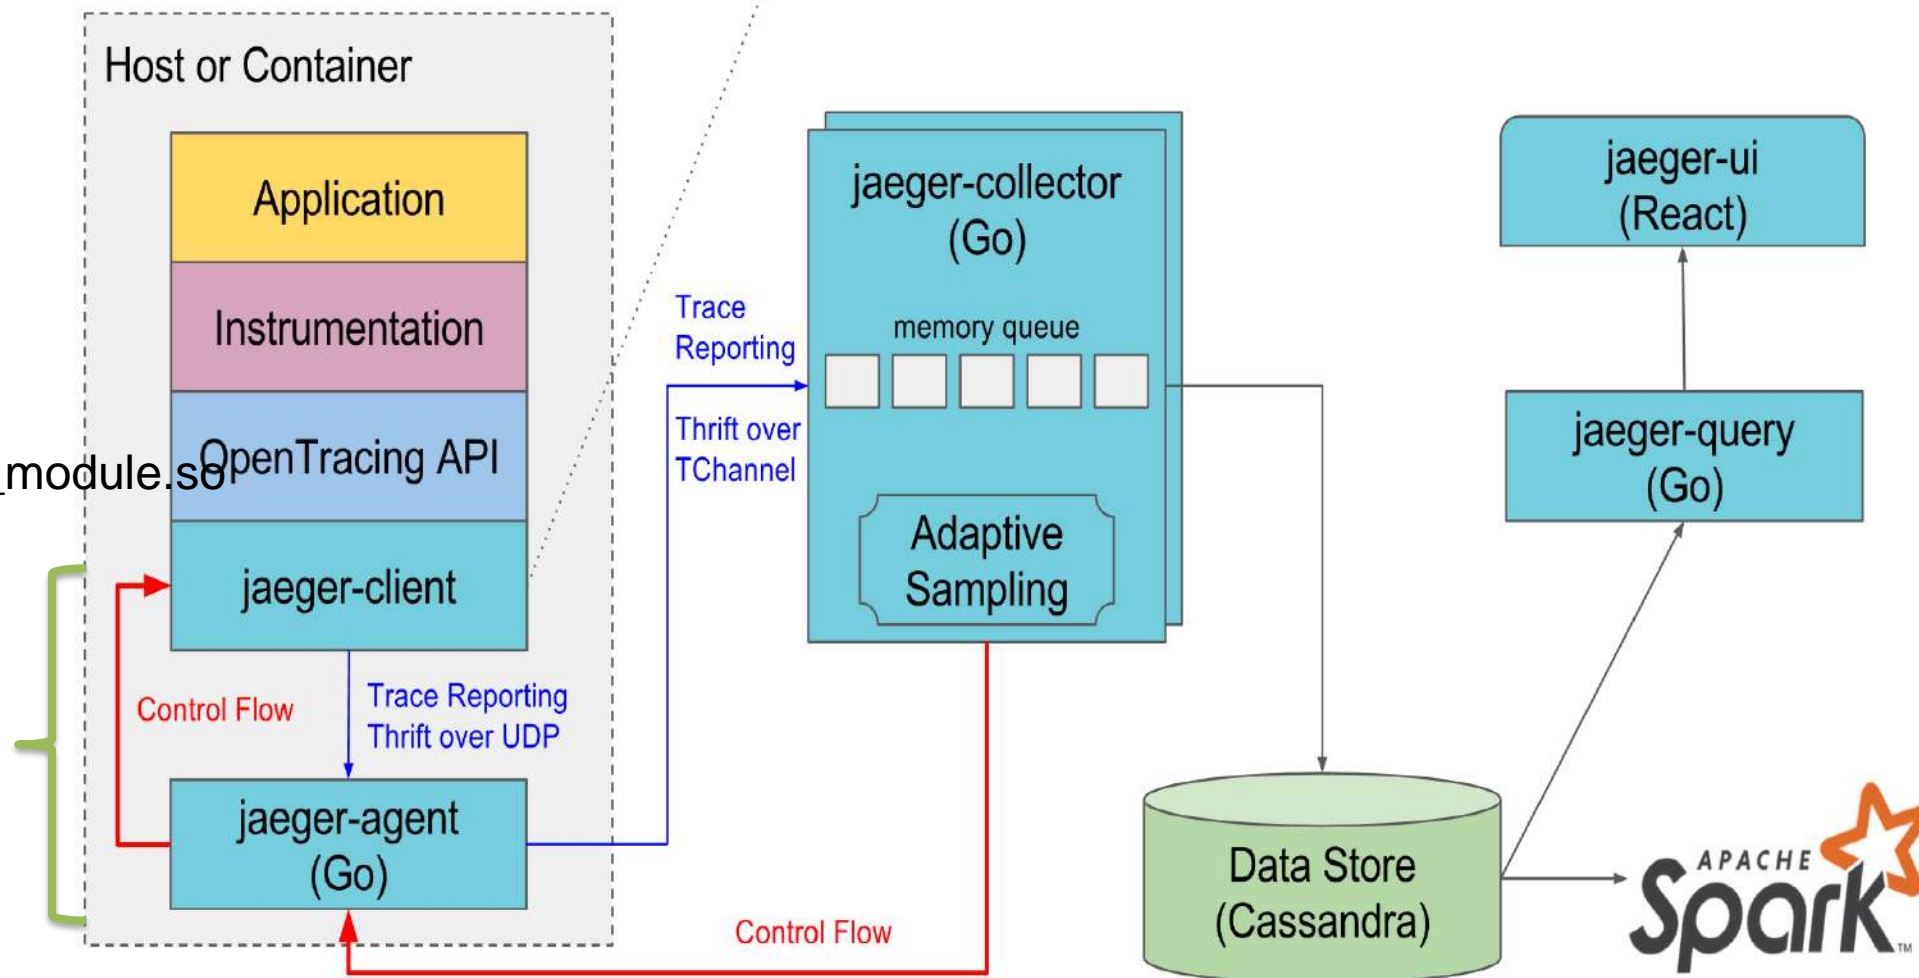
JAEGER



OPENTRACING

ngx_http_opentracing_module.so

Jaeger plugin
vendor tracer




```

# Load the OpenTracing dynamic module.
load_module modules/nginx_http_opentracing_module.so;

http {
    # Load a vendor tracer
    opentracing_load_tracer /usr/local/lib/libjaegertracing_plugin.so /etc/jaeger-nginx-config.json;

    # or
    # opentracing_load_tracer /usr/local/lib/liblightstep_tracer_plugin.so /path/to/config;
    # or
    # opentracing_load_tracer /usr/local/lib/libzipkin_opentracing_plugin.so /path/to/config;
    # or
    # opentracing_load_tracer /usr/local/lib/libdd_opentracing_plugin.so /path/to/config;

    # Enable tracing for all requests.
    opentracing on;

    # Optionally, set additional tags.
    opentracing_tag http_user_agent $http_user_agent;

    upstream backend {
        server app-service:9001;
    }

    location ~ {
        # The operation name used for spans defaults to the name of the location
        # block, but you can use this directive to customize it.
        opentracing_operation_name $uri;

        # Propagate the active span context upstream, so that the trace can be
        # continued by the backend.
        # See http://opentracing.io/documentation/pages/api/cross-process-tracing.html
        opentracing_propagate_context;

        proxy_pass http://backend;
    }
}

```

```

# Load a vendor tracer
opentracing_load_tracer /usr/local/lib/libjaegertracing_plugin.so
                        /etc/jaeger/jaeger-config.json;
#opentracing_load_tracer /usr/local/lib/libzipkin_opentracing_plugin.so
#                        /etc/zipkin/zipkin-config.json;

# Enable tracing for all requests
opentracing on;

# Set additional tags that capture the value of NGINX Plus variables
opentracing_tag bytes_sent $bytes_sent;
opentracing_tag http_user_agent $http_user_agent;
opentracing_tag request_time $request_time;
opentracing_tag upstream_addr $upstream_addr;
opentracing_tag upstream_bytes_received $upstream_bytes_received;
opentracing_tag upstream_cache_status $upstream_cache_status;
opentracing_tag upstream_connect_time $upstream_connect_time;
opentracing_tag upstream_header_time $upstream_header_time;
opentracing_tag upstream_queue_time $upstream_queue_time;
opentracing_tag upstream_response_time $upstream_response_time;

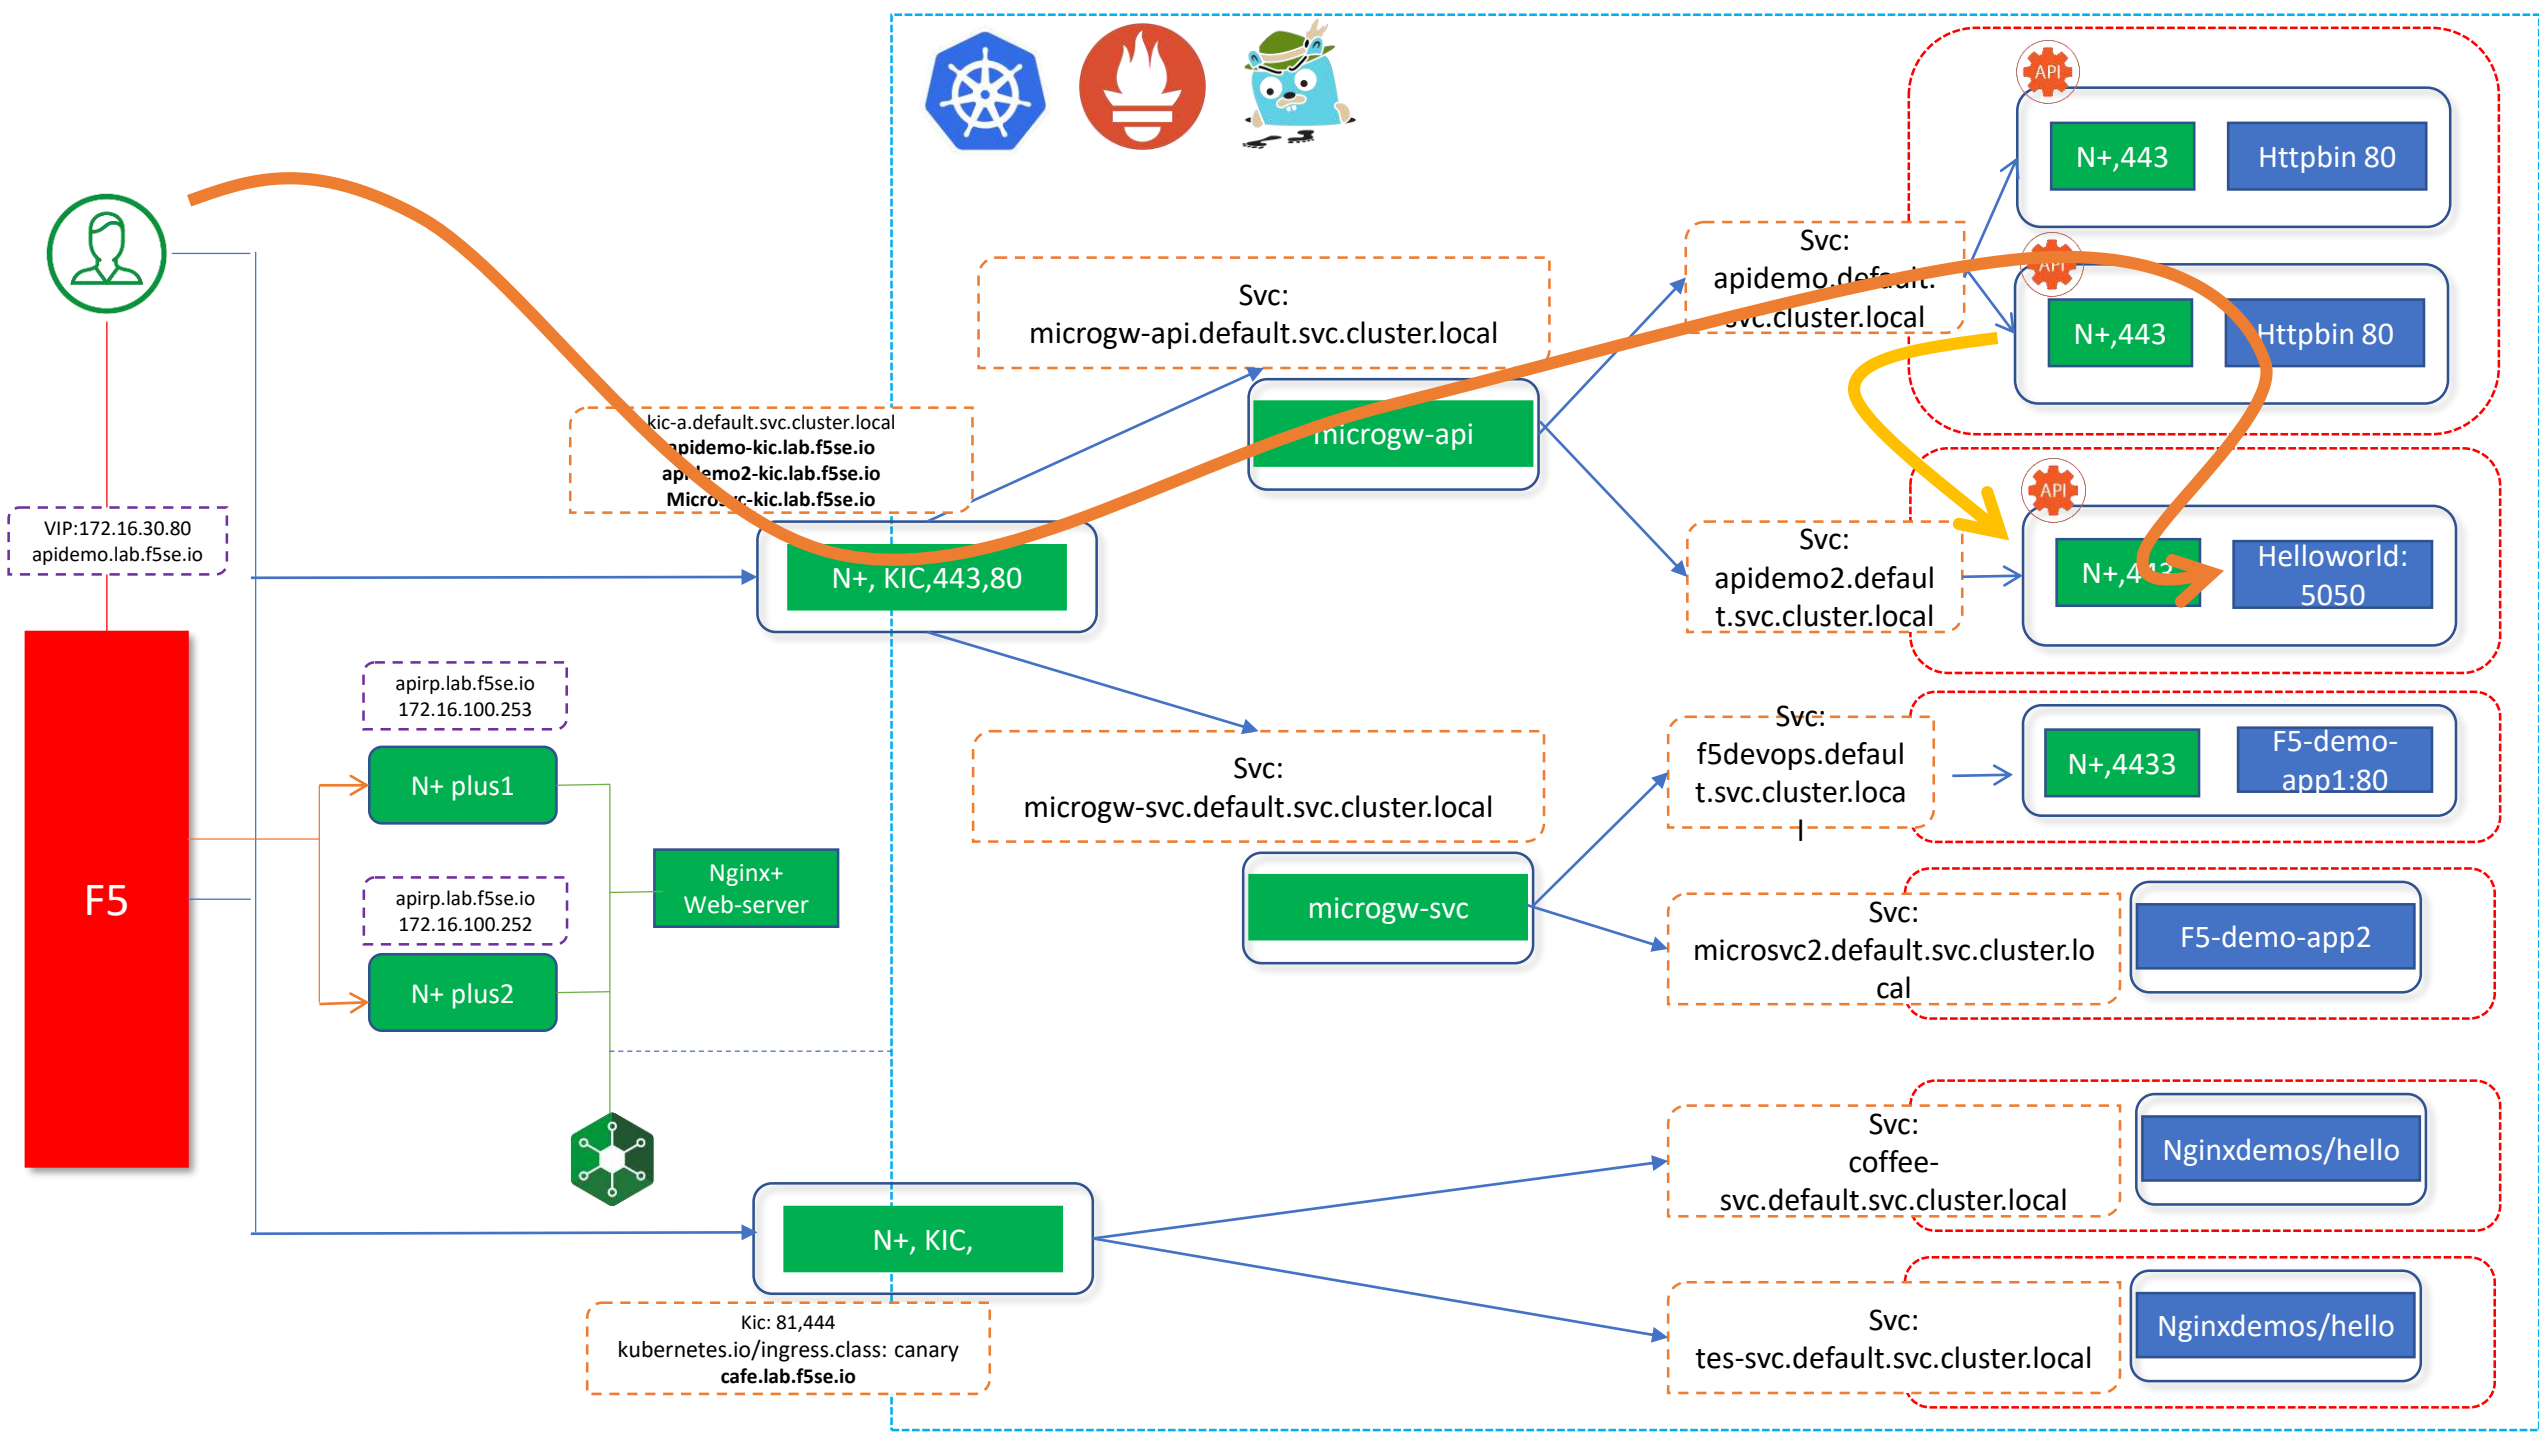
server {
    listen 9001;

    location / {
        # The operation name used for OpenTracing Spans defaults to the name of the
        # 'location' block, but uncomment this directive to customize it.
        #opentracing_operation_name $uri;

        # Propagate the active Span context upstream, so that the trace can be
        # continued by the backend.
        opentracing_propagate_context;

        # Make sure that your Ruby app is listening on port 4567
        proxy_pass http://127.0.0.1:4567;
    }
}

```





```
kubectl get cm -n nginx-ingress nginx-config -o yaml
```

观察jaeger配置

模拟一个问题

kubectl get pod,svc -o wide

[访问 https://apidemo-kic.lab.f5se.io/gotohelloworld/d](https://apidemo-kic.lab.f5se.io/gotohelloworld/d)
得到正常返回

删除并重建 hello-world service:

kubectl delete -f hello-world-sidecar-svc-deploy.yaml

kubectl create -f hello-world-sidecar-svc-deploy.yaml

再次访问:

[访问 https://apidemo-kic.lab.f5se.io/gotohelloworld/d](https://apidemo-kic.lab.f5se.io/gotohelloworld/d)

是否可以访问？为什么？查看jaeger可以发现什么？

<http://jaeger.lab.f5se.io:32686/>

原因是什么？怎么解决？



```
[root@k8s-master nginx-sidecar]# kubectl get pod,svc -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
pod/helloworld-sidecar-6b76979dd4-p7jc5	2/2	Running	0	22m	10.244.1.216	k8s-node1
pod/httpbin-sidecar-7ffcfd7b49-w58zd	2/2	Running	0	16m	10.244.1.217	k8s-node1
pod/jaeger-65fcbf66db-9gt2q	1/1	Running	1	9d	10.244.0.227	k8s-master
pod/microgw-api-nginx-tracing-57bf585bf4-bg56w	1/1	Running	1	9d	10.244.0.224	k8s-master
pod/tools	0/1	ImagePullBackOff	0	1h	10.244.0.229	k8s-master

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
service/apidemo	ClusterIP	10.250.0.187	<none>	443/TCP,80/TCP	16m	app=httpbin-sidecar
service/apidemo2	ClusterIP	10.250.0.157	<none>	443/TCP,80/TCP	22m	app=helloworld-sidecar
service/jaeger-agent	ClusterIP	10.250.0.122	<none>	5775/UDP,6831/UDP,6832/UDP,5778/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger
service/jaeger-collector	ClusterIP	10.250.0.246	<none>	14267/TCP,14268/TCP,9411/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger
service/jaeger-query	NodePort	10.250.0.139	<none>	80:32686/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger
service/kubernetes	ClusterIP	10.250.0.1	<none>	443/TCP	1y	<none>
service/microgw-api	NodePort	10.250.0.23	<none>	443:32610/TCP,80:32014/TCP	9d	app=microgw-api-nginx-tracing
service/nginx-deploy-svc	ClusterIP	10.250.0.117	<none>	80/TCP	100d	app=nginx-deploy
service/zipkin	ClusterIP	None	<none>	9411/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger

正常访问



⚠ Not secure | apidemo-kic.lab.f5se.io/gotohelloworld/d

Apps k8s-Container ServiceMesh Indigo Archi ELK DevOps Openstack F5 company

/:path1/:path2 - Hello to gotohelloworld/d ! Host:helloworld-sidecar-6b76979dd4-p7jc5/10.244.1.216

nginx-ingress-api: /gotohelloworld/d 5e15d03

8 Spans

microgw (2)

nginx-ingress-api (2)

sidecar-app-helloworld (2)

sidecar-app-httpbin (2)

删除并重建hello world service

```
kubectl delete -f hello-world-sidecar-svc-deploy.yaml
```

```
kubectl create -f hello-world-sidecar-svc-deploy.yaml
```

```
[root@k8s-master nginx-sidecar]# kubectl delete -f hello-world-sidecar-svc-deploy.yaml
service "apidemo2" deleted
deployment.extensions "helloworld-sidecar" deleted
[root@k8s-master nginx-sidecar]# kubectl create -f hello-world-sidecar-svc-deploy.yaml
service "apidemo2" created
deployment.extensions "helloworld-sidecar" created
[root@k8s-master nginx-sidecar]# kubectl get pod,svc -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
pod/helloworld-sidecar-6b76979dd4-cf6wp	2/2	Running	0	15s	10.244.0.232	k8s-master
pod/helloworld-sidecar-6b76979dd4-p7jc5	2/2	Terminating	0	25m	10.244.1.216	k8s-node1
pod/httpbin-sidecar-7ffcfd7b49-w58zd	2/2	Running	0	20m	10.244.1.217	k8s-node1
pod/jaeger-65fcbf66db-9gt2q	1/1	Running	1	9d	10.244.0.227	k8s-master
pod/microgw-api-nginx-tracing-57bf585bf4-bg56w	1/1	Running	1	9d	10.244.0.224	k8s-master
pod/tools	0/1	ImagePullBackOff	0	1h	10.244.0.229	k8s-master

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
service/apidemo	ClusterIP	10.250.0.187	<none>	443/TCP,80/TCP	20m	app=httpbin-sidecar
service/apidemo2	ClusterIP	10.250.0.107	<none>	443/TCP,80/TCP	15s	app=helloworld-sidecar
service/jaeger-agent	ClusterIP	10.250.0.122	<none>	5775/UDP,6831/UDP,6832/UDP,5778/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger
service/jaeger-collector	ClusterIP	10.250.0.246	<none>	14267/TCP,14268/TCP,9411/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger
service/jaeger-query	NodePort	10.250.0.139	<none>	80:32686/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger
service/kubernetes	ClusterIP	10.250.0.1	<none>	443/TCP	1y	<none>
service/microgw-api	NodePort	10.250.0.23	<none>	443:32610/TCP,80:32014/TCP	9d	app=microgw-api-nginx-tracing
service/nginx-deploy-svc	ClusterIP	10.250.0.117	<none>	80/TCP	100d	app=nginx-deploy
service/zipkin	ClusterIP	None	<none>	9411/TCP	9d	app.kubernetes.io/component=all-in-one,app.kubernetes.io/name=jaeger



apidemo-kic.lab.f5se.io/gotohelloworld/d

fesh Indigo Archi ELK DevOps Openstack F5 company Ansible SE-LAB linux 安全 生活 Python AWS/Azure

An error occurred.

Sorry, the page you are looking for is currently unavailable.
Please try again later.

If you are the system administrator of this resource then you should check the error log for details.

Faithfully yours, nginx.

nginx-ingress-api: /gotohelloworld/d cf633dd

7 Spans 6 Errors

microgw (2) nginx-ingress-api (2) sidecar-app-httpbin (3)

microgw /

sidecar-app-httpbin /50x.html

sidecar-app-httpbin /gotohelloworld

sidecar-app-httpbin /50x.html

/50x.html

Tags

peer.address	"10.244.0.224:57558"
http.method	"GET"
http.url	"https://apidemo-kic.lab.f5se.io/gotohelloworld/d"
http.host	"apidemo-kic.lab.f5se.io"
bytes_sent	"0"
http_user_agent	"Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Ge"
request_time	"12.000"
upstream_addr	"10.250.0.157:443"
upstream_bytes_received	"0"
upstream_cache_status	" "
upstream_connect_time	"."
upstream_header_time	"."
upstream_queue_time	"0.000"
upstream_response_time	"12.001"
http.status_code	502
http.status_desc	" "



原因

```
proxy_pass https://apidemo2.default.svc.cluster.local;
```



怎么解决？

```
resolver 10.250.0.53 valid=3s;  
upstream apidemo2backends {  
    zone apidemo2backends 32k;  
    server apidemo2.default.svc.cluster.local service=_https._tcp resolve;  
}
```

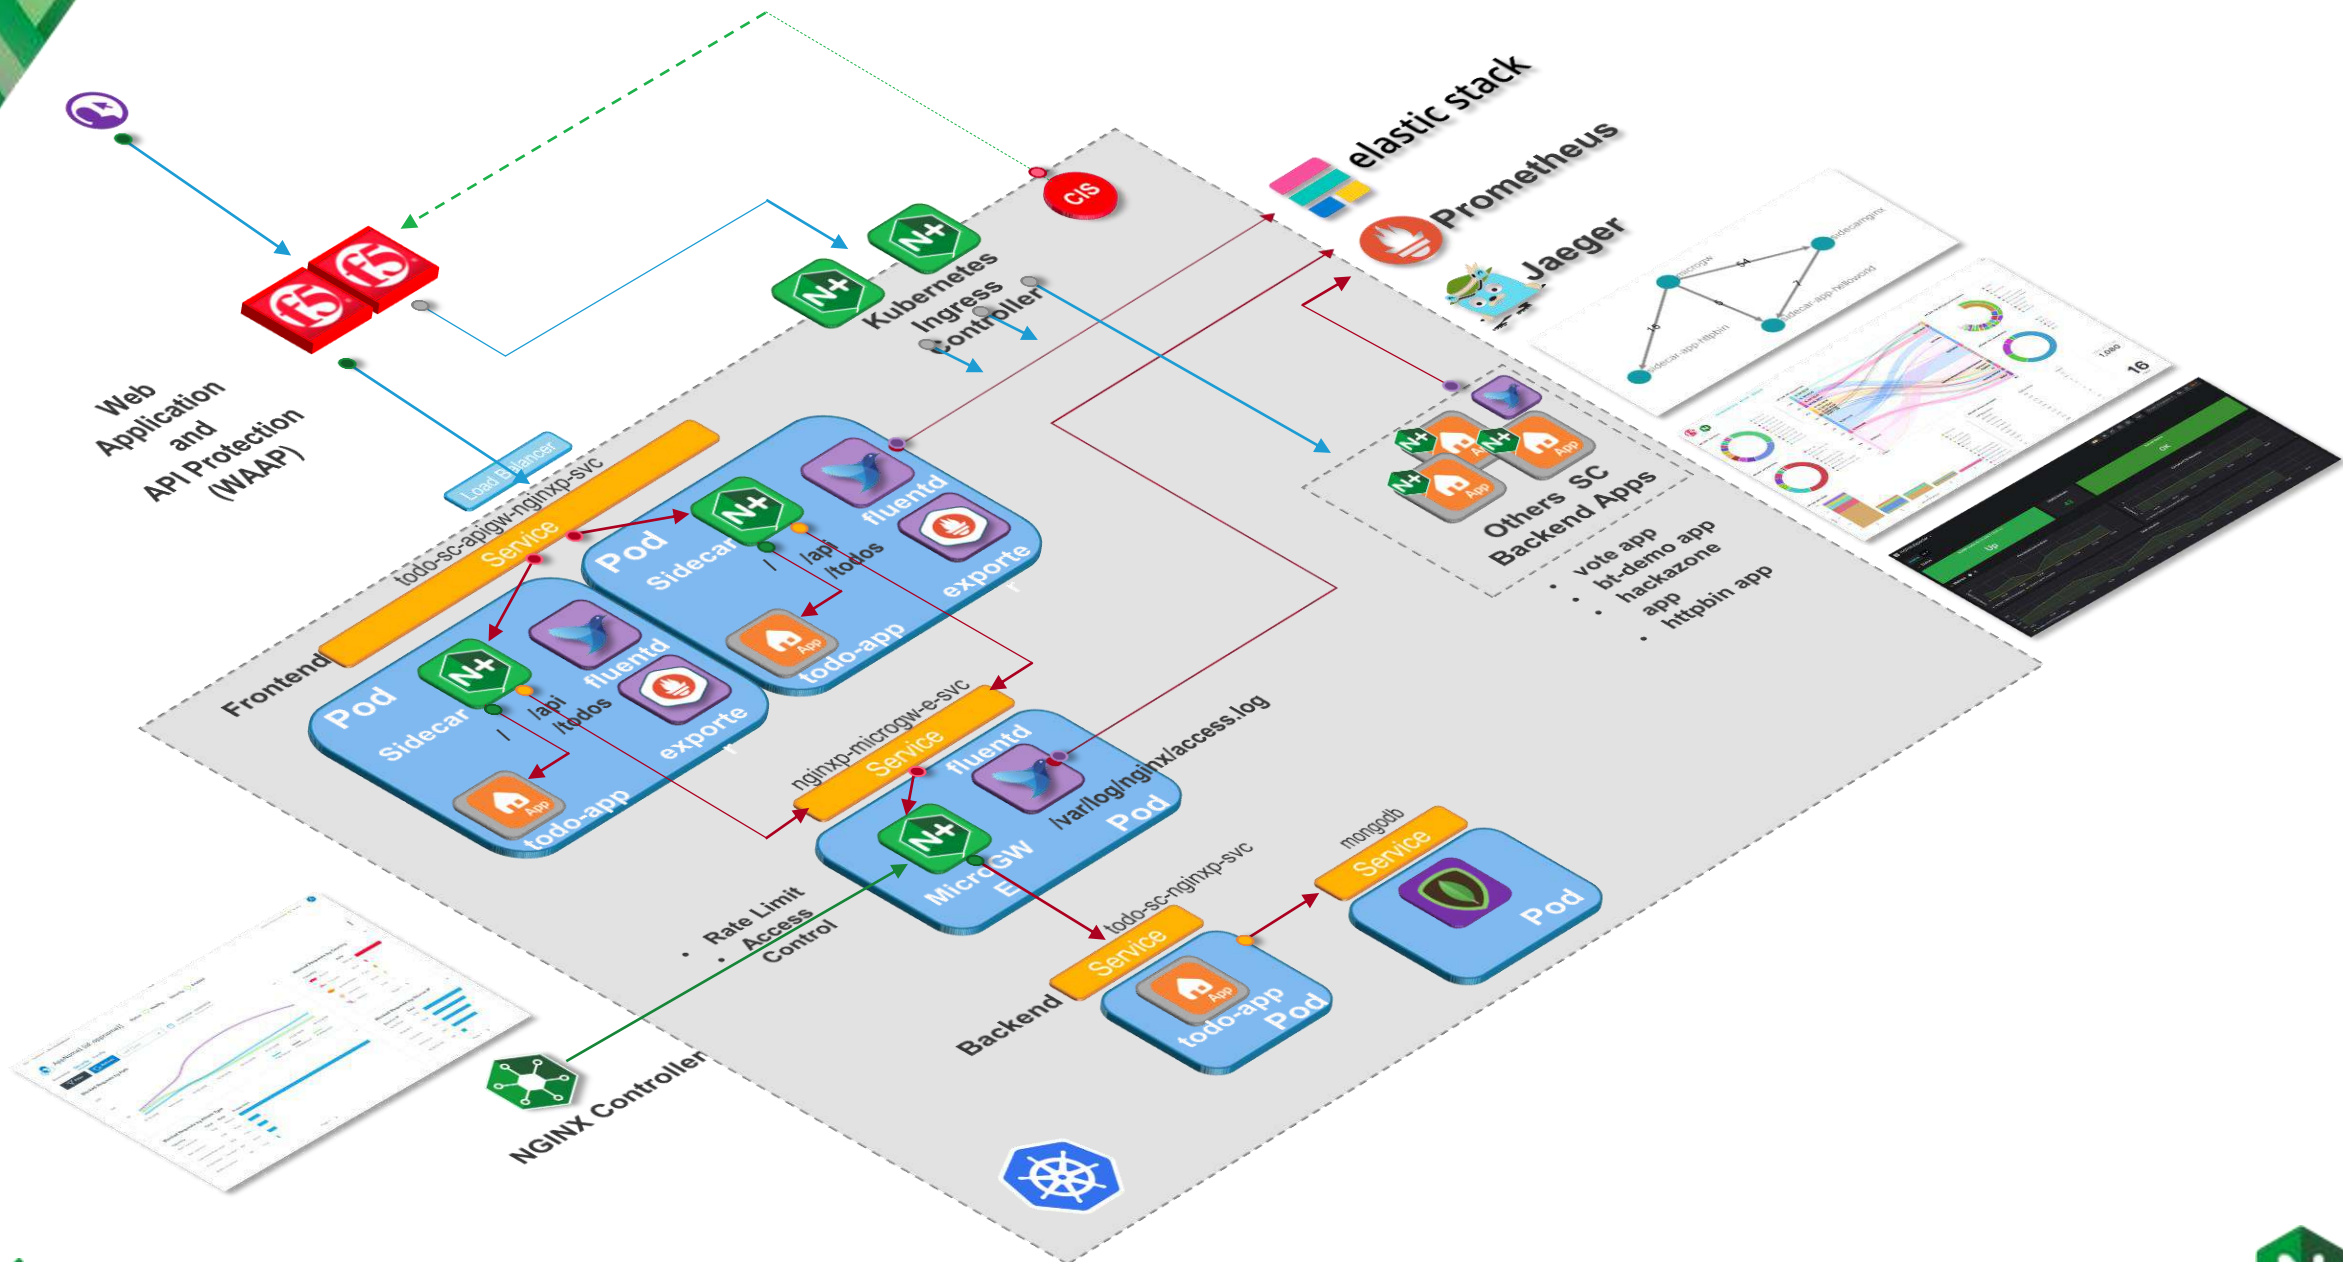


留个问题

Upstream的上游是多个公有云中的服务，例如 `svc.aws.com`, `svc.azure.com`, `svc.ali.com`
且这些服务对应的IP可能会发生变化

NGINX作为反向代理，应如何设置upstream才能确保随时可以正常访问这些业务？

<https://www.myf5.net/post/2791.htm>





Thank You

关注我们

F5 官方微信
(新闻, 技术文章)



F5社区
(答疑, 吐槽, 分享, 互动)



加入F5社区：关注“F5社区”微信公众号，注册成为社员，随时参加meet up活动，代码共享，讨论，答疑等。只要你有想法，有创造，那么就快来大展身手吧，让我们在社区里尽情分享，交流，吐槽和互动，在这个自由的国度里，发现闪亮的自己。让我们一起来见证“一群有才能的人在一起做有梦想的事！”

NGINX技术群



操作步骤：

1. 扫描二维码并在“入群信息”栏填写姓名
2. 点击下方“我要入群”
3. 长按识别二维码进入群聊